

FairClaim

Working name — AI-drafted demand letters and consumer dispute resolution

Execution & Business Plan · Build with Gemini XPRIZE · Professional Services category

Prepared July 1, 2026 · Submission deadline August 17, 2026, 1:00 PM PT · 47 days remaining

1. Executive Summary

FairClaim turns a consumer's dispute into a professionally formatted, statute-aware demand letter in about ten minutes, for a flat fee of \$39. The target buyer is someone who has just been wronged in a way with a clear dollar value attached: a landlord withholding a security deposit, a contractor who took a deposit and disappeared, an airline refusing a refund it legally owes, a gym that keeps billing after cancellation. These people will not pay a lawyer \$300–500 to recover \$800, and most never pursue the claim at all. A credible demand letter on their behalf resolves a meaningful share of these disputes without anyone going to court, and for the rest it is the required first step toward small claims filing.

The business is built to win under this competition's specific judging criteria, which weight business viability, AI-native operations, and category impact equally. FairClaim charges from day one, so revenue evidence accumulates across the whole window. Its operations are genuinely AI-native: a Gemini-powered agent pipeline handles intake, jurisdiction research, drafting, quality review, and escalation guidance with no human in the loop for the standard product. And its impact story maps directly onto the Professional Services category framing used in XPRIZE's own launch messaging about closing gaps in access to legal services.

The plan below allocates the 47 remaining days roughly as follows: three days of demand validation with manual fulfillment, eleven days of product build, and thirty days where the overwhelming majority of effort goes to distribution, revenue, and testimonial collection, with the final three days reserved for assembling the submission package. The single most important strategic commitment in this document is that build time is capped: after July 14, feature work only happens in service of a demonstrated sales blocker.

2. The Problem and the Customer

Small-dollar civil disputes are the largest underserved segment of the legal market. The economics are structurally broken: when the amount in controversy is \$500–5,000, no lawyer can profitably take the case, so the injured party's real options are writing an angry email themselves (usually ignored), navigating small claims court alone (intimidating, slow, unfamiliar), or walking away (what most people do). The gap between "angry email" and "lawyer" is exactly where a \$39 product lives.

The buyer profile that matters for this competition is defined less by demographics than by a trigger event. Someone whose landlord just kept a \$1,400 deposit is searching for answers within hours, is emotionally motivated, and makes a purchase decision in minutes rather than weeks. That trigger-driven urgency is what makes a 47-day revenue window feasible: there is no long consideration cycle to wait out. The four launch verticals, chosen for volume, statute clarity, and reachability, are security deposit disputes, contractor and home-services disputes, airline and travel refunds, and subscription or billing disputes.

Each has well-defined statutory hooks (state deposit-return deadlines and penalty multipliers, state consumer protection acts, DOT refund rules, auto-renewal laws) that make the letters concretely stronger than anything the customer would write alone.

Where these buyers can be reached quickly matters as much as who they are. They self-identify in public, searchable places at the moment of need: tenant-rights and legal-advice adjacent forums, city subreddits, consumer complaint communities, and Facebook groups for renters and homeowners. Section 6 treats these channels in detail.

3. Product Definition

3.1 Core product

The customer completes a structured intake (dispute type, state, timeline, amounts, uploads of any supporting documents such as a lease or receipts). The agent pipeline produces a formal demand letter citing the relevant statutes for their jurisdiction, a plain-English cover note explaining what the letter says and why those citations matter, and a recommended delivery method (certified mail language is included in the letter). Delivery is a polished PDF plus editable text, emailed and available in a simple account page. Target time from payment to delivery: under ten minutes.

3.2 Pricing and packaging

Tier	Price	What's included
Demand Letter	\$39	Statute-cited demand letter, plain-English cover note, delivery instructions, one free revision.
Letter + Escalation Pack	\$59	Everything above, plus a follow-up letter template, a small-claims filing roadmap specific to the customer's county where determinable, and an evidence-organization checklist.
Response Review (add-on)	\$19	Customer forwards the other party's reply; the agent analyzes it and drafts the recommended next move.

The \$39 anchor is deliberately impulse-priced against the customer's recovery amount: on a typical \$1,000+ dispute it reads as trivially cheap insurance. The escalation pack exists mainly to raise average order value; the response-review add-on creates a legitimate second transaction and a sustainability story for judges. Refund policy is unconditional within 7 days, which removes purchase risk and generates goodwill in the communities where reputation drives referrals.

3.3 Explicitly out of scope for the competition window

No attorney network, no marketplace, no B2B offering, no mobile app, no multi-language support, and no verticals beyond the four listed. Each of these is a plausible future direction and each is a trap in a 47-day window. The submission can mention them as roadmap; none should consume a build hour before August 17.

4. Fit to the Judging Criteria

Submissions are judged on three equally weighted criteria. The table below is the strategic spine of this plan; every week's work should strengthen at least one row, and the submission write-up should be organized around them.

Criterion	FairClaim's evidence	What to capture along the way
Business viability	Paying customers from week one; flat-fee pricing with ~90% gross margin; clear unit economics (Gemini + infra cost well under \$1 per order); repeat-purchase mechanics via response review.	Stripe revenue dashboard screenshots over time, order counts by week, conversion rates by channel, average order value.
AI-native operations	The entire service is delivered by an agent pipeline: intake validation, jurisdiction research, drafting, adversarial QA review, and escalation guidance run on Gemini with zero marginal human labor. The business could not exist at this price point without AI.	Architecture diagram, pipeline logs/metrics (orders processed, median delivery time), a short screen recording of an order flowing through untouched.
Category impact	Access-to-justice framing: customers recover money they would otherwise have abandoned, at 1/10th the cost of the cheapest human alternative. Directly echoes XPRIZE's launch language about closing gaps in legal services access.	Customer outcome stories (with permission): dollars recovered, disputes resolved without court. A running "total recovered for customers" counter is the single most powerful impact metric.

Note the asymmetry: viability and impact evidence can only be gathered while the competition is running. This is the deeper reason the plan caps build time — every extra week of building is a week of evidence that cannot be recovered later.

5. The 47-Day Timeline

Dates run from Wednesday, July 1 through the August 17 deadline. The plan front-loads validation so that if the core assumption fails (people will pay \$39 for this), there is still time to pivot to the freelancer contract-review concept documented as the fallback in Section 9.

Phase	Dates	Objective and key activities
Phase 0: Validate	Jul 1–3	Prove someone will pay before writing product code. Register for the competition on Devpost and geminixprize.com. Set up a one-page site with a Stripe payment link and an intake form (form tool, no custom code). Post helpfully in 3–4 target communities. Fulfill the first orders manually using Gemini in AI Studio. Success gate: 3+ paid orders or equivalently strong signal by end of July 3.
Phase 1: Build MVP	Jul 4–14	Build the real product per the Technical Specification document: intake flow, agent pipeline, payment, PDF delivery, order dashboard. Continue fulfilling orders manually throughout — manual orders are both revenue and training data for prompt refinement. Ship to production on Cloud Run by July 14 even if imperfect.

Phase	Dates	Objective and key activities
Phase 2: Launch & first traction	Jul 15–21	Switch fulfillment to the automated pipeline with human review of every letter for the first week (review, not rewrite). Launch publicly: Product Hunt or equivalent, the full community list, and begin the content cadence in Section 6. Target: 25 cumulative orders by July 21.
Phase 3: Distribution grind	Jul 22–Aug 10	Nearly 100% of effort on sales and customer outcomes. Daily community engagement, respond to every customer, collect testimonials and recovery outcomes, publish 2–3 vertical-specific landing pages (e.g., "[State] security deposit demand letter") for search and shareability. Feature work only to unblock demonstrated sales problems. Target: 100+ cumulative orders, first documented customer recoveries.
Phase 4: Submission	Aug 11–17	Assemble the package per Section 10: 3-minute demo video, write-up organized around the three judging criteria, metrics screenshots, customer stories, architecture diagram. Submit by August 15 to leave a two-day buffer against the hard August 17, 1:00 PM PT deadline.

Two rules govern the timeline. First, the Phase 0 gate is real: if nobody pays by July 3 despite genuine reach into the target communities, the concept pivots rather than proceeding on hope. Second, the August 15 internal deadline is treated as the true deadline; Devpost submissions have a long history of last-hour upload failures and the rules give no grace period.

6. Go-to-Market Playbook

6.1 Channel strategy

With no time for SEO to mature and no budget assumed for paid acquisition, the strategy is community-led: be genuinely useful in the places where people describe these disputes, in a way that makes the product the obvious next step. The core motion is answering real posts with real substance (what the relevant statute says, what a demand letter should contain, what the small-claims path looks like) and offering the product as the convenient version of the advice just given. Most communities ban naked self-promotion but tolerate and even welcome expert-flavored help with a disclosed affiliation; read each community's rules first and when in doubt, give the help and let the profile link do the selling.

Channel	Motion	Cadence
Reddit (tenant, consumer-rights, city, and legal-advice-adjacent subreddits)	Substantive answers on live dispute threads; one dedicated launch/feedback post in entrepreneurship and side-project communities.	45–60 min daily
Facebook groups (renters, landlord-tenant, homeowner, travel complaint groups)	Same helpful-answer motion; groups skew older and convert well on done-for-you services.	30 min daily

Channel	Motion	Cadence
Vertical landing pages	Pages per dispute type and major state ("California security deposit demand letter") that answer the query fully, then offer the product. Shareable in threads as a resource, which sidesteps most self-promotion rules.	2–3 pages/week in Phase 3
Short-form video (optional, if energy allows)	60-second explainers: "Your landlord has X days to return your deposit in [state] — here's the letter that gets it back." Consumer-rights content performs reliably on TikTok/Reels/Shorts.	2–3/week
Launch platforms	Product Hunt / Hacker News style launch during Phase 2 for an early spike, backlinks, and social proof; not a durable channel.	Once

6.2 Positioning and message

The message that converts is specific, monetary, and deadline-shaped: "Your landlord had 21 days. It's been 40. For \$39, send a demand letter citing the exact statute they're violating — most deposits come back without court." Every landing page and post lead should name the statute clock, the dollar stakes, and the ten-minute turnaround. Avoid the word "AI" in customer-facing copy except where it explains speed and price; customers are buying an outcome, not a technology. (The competition submission, by contrast, should foreground the AI architecture heavily — different audiences.)

6.3 The testimonial engine

Every delivered order triggers a follow-up sequence at day 7 and day 21 asking one question: "Did they respond? Did you get your money?" Positive outcomes convert into (a) permissioned testimonials, (b) entries in the running total-recovered counter, and (c) anonymized before/after stories for community posts, which are themselves the best-performing content in this niche. This loop is the bridge between the go-to-market and the impact evidence the judges need to see, and it should be built in Phase 1, not bolted on later.

7. Revenue Model and Targets

Unit economics: at \$39–59 per order, the marginal cost per order (Gemini API calls across the pipeline, PDF generation, hosting, Stripe fees) is roughly \$2–3, of which Stripe is the largest component. Gross margin is approximately 93%. There is no marginal labor. The \$300 in Cloud credits and \$100 in AI credits granted at registration likely cover most infrastructure through the entire window.

Scenario	Orders by Aug 15	Revenue	What it demonstrates
Floor	40	~\$1,800	Real strangers pay; concept validated; competitive for category and runner-up prizes if impact stories are strong.
Base	100	~\$4,500	Repeatable channel economics; week-over-week growth curve; strong runner-up and category contention.

Scenario	Orders by Aug 15	Revenue	What it demonstrates
Stretch	250+	~\$11,000+	Escape velocity; documented recoveries exceeding \$50k returned to customers; finalist-tier evidence.

The judges are evaluating trajectory and credibility, not absolute scale — no 90-day business is large. A clean week-over-week growth chart with real customer outcomes beats a larger but flat or astroturfed number. All revenue flows through Stripe from day one so the entire history is independently screenshot-able.

8. Legal Positioning and Trust

The central legal consideration is the unauthorized practice of law (UPL). FairClaim's posture is self-help document preparation: the customer is representing themselves, and the product prepares documents at their direction — the same posture LegalZoom, Rocket Lawyer, and similar services have operated under for two decades. Concretely this means: prominent "not a law firm, not legal advice" disclaimers at intake, in the cover note, and in the footer; letters written in the customer's own voice and signed by the customer, never by FairClaim; no individualized advice about whether to sue or settle (the escalation roadmap presents procedural information, not recommendations); and terms of service stating the customer is responsible for reviewing before sending. This document is planning input, not legal advice — before scaling past the competition window, a one-hour consult with a lawyer familiar with legal-tech UPL questions is a worthwhile \$300.

Data handling is the second trust pillar: intake documents (leases, receipts) are personal. Commitments worth making publicly and cheaply honoring: documents are used only to produce the customer's letter, auto-deleted after 30 days, never used to train models, and processed on Google Cloud infrastructure. The technical spec covers implementation.

9. Risks and Mitigations

Risk	Likelihood	Mitigation
Phase 0 validation fails — community posts get engagement but nobody pays.	Moderate	Pre-planned pivot to the adjacent validated concept (freelancer contract review, same pipeline architecture, ~80% code reuse). Decision point is hard-scheduled for July 3.
Letter quality embarrassment — a wrong statute citation surfaces publicly.	Moderate	Adversarial QA agent checks every citation against a curated statute reference set rather than trusting model recall; human review of 100% of letters in the first automated week; free revision policy converts errors into service recoveries.
Community moderation removes posts / bans account.	High for at least one community	Diversify across 8–10 communities; lead with substance; disclose affiliation; landing pages as linkable resources rather than direct pitches.

Risk	Likelihood	Mitigation
Solo-founder time crunch — build eats the selling window.	High (the default failure mode)	Hard build cap on July 14; manual fulfillment keeps revenue flowing during the build; scope in Section 3.3 is contractual with yourself.
Platform/verification issues at submission (Devpost upload, eligibility paperwork, Gemini-usage verification).	Low	Register and read official rules in full during Phase 0; submit August 15; keep pipeline logs demonstrating Gemini usage throughout.

10. Submission Package Checklist

Everything below is assembled in Phase 4, but the evidence it depends on is captured continuously from July 1. Screenshot the Stripe dashboard and order metrics weekly; future-you will want the time series.

Item	Notes
Demo video (~3 min)	Structure: 30s problem (a real deposit-dispute story), 90s product walkthrough showing an order flowing through the agent pipeline end-to-end, 60s traction and impact (revenue chart, total recovered counter, testimonial). Record with real UI, not slides.
Written submission	Organized explicitly under the three judging criteria headings so screeners can score without hunting. Include the architecture diagram from the technical spec.
Traction evidence	Stripe screenshots (weekly series), order count chart, channel conversion table, median delivery time metric.
Impact evidence	Total recovered for customers; 3–5 permissioned customer stories with dollar outcomes; cost comparison vs. attorney demand letter (\$39 vs. \$300–500).
AI-native evidence	Pipeline architecture diagram, anonymized pipeline logs, screen recording of zero-touch fulfillment, note on Gemini products used (Gemini API, AI Studio, Cloud Run).
Compliance items	Eligibility confirmation, hosted URL, repo access if requested, completed Devpost form fields — re-read the official rules checklist at geminixprize.com/rules before submitting.

Final note on positioning the write-up: the strongest available narrative is that this business could not exist without AI — at \$39, human fulfillment is impossible, and before this technology the customer's money simply went unrecovered. That sentence connects all three judging criteria and should appear, in some form, in the first paragraph of the submission.